

Task 2 Data Exploration with Python

a) Exploratory Data Analysis in Python

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("GHG_Data.csv")

# Overview of the dataset
print(df.head())
print(df.info())
print(df.describe())
```

	Year	Ontario	GHG ID	Facility Owner \
0	2010		1001	ADM Agri-Industries
1	2010		1002	Air Products Canada Ltd
2	2010		1003	Algonquin Power Energy From Waste Inc.
3	2010		1005	ArcelorMittal Dofasco Inc.
4	2010		1006	Atlantic Packaging Products Ltd.

	Facility Name	Facility City \
0	ADM Windsor	Windsor
1	Corunna Hydrogen Facility	Corunna
2	Algonquin Power Energy from Waste Inc.	Brampton
3	Dofasco Hamilton	Hamilton
4	111 Progress	Scarborough

	Facility	Primary NAICS Code \
0		311224
1		325120
2		562210
3		331110
4		332120

b) Summary Statistics

```
In [15]: #import pip
await pip.install('seaborn')
```

```
In [14]: import seaborn as sns
```

```
# Display the first few rows of the dataset
print(df.head())

# Get the basic statistics of numerical columns
print(df.describe())

# Get information about the data types and missing values
print(df.info())
```

	Year	Ontario GHG ID	Facility Owner \
0	2010	1001	ADM Agri-Industries
1	2010	1002	Air Products Canada Ltd
2	2010	1003	Algonquin Power Energy From Waste Inc.
3	2010	1005	ArcelorMittal Dofasco Inc.
4	2010	1006	Atlantic Packaging Products Ltd.

	Facility Name	Facility City \
0	ADM Windsor	Windsor
1	Corunna Hydrogen Facility	Corunna
2	Algonquin Power Energy from Waste Inc.	Brampton
3	Dofasco Hamilton	Hamilton
4	111 Progress	Scarborough


	Facility Primary NAICS Code \
0	311224
1	325120
2	562210
3	331110
4	332120

```
In [16]: # Plot a histogram for the "Total CO2e from all sources in CO2e (t)" column
plt.figure(figsize=(8, 6))
sns.histplot(df['Total CO2e from all sources in CO2e (t)'], bins=30, kde=True)
plt.xlabel('Total CO2e (t)')
plt.ylabel('Frequency')
plt.title('Distribution of Total CO2e')
plt.show()

# Scatter plot between "Carbon dioxide (CO2) from non-biomass in CO2e (t)" and
plt.figure(figsize=(8, 6))
sns.scatterplot(x='Carbon dioxide (CO2) from non-biomass in CO2e (t)', y='Total CO2e from all sources in CO2e (t)')
plt.xlabel('CO2 from non-biomass (t)')
plt.ylabel('Total CO2e (t)')
plt.title('Scatter Plot of CO2 from non-biomass vs. Total CO2e')
plt.show()

# Box plot for the "Facility City" vs. "Total CO2e"
plt.figure(figsize=(10, 6))
sns.boxplot(x='Facility City', y='Total CO2e from all sources in CO2e (t)', data=df)
plt.xticks(rotation=45, ha='right')
plt.xlabel('Facility City')
plt.ylabel('Total CO2e (t)')
plt.title('Box Plot of Total CO2e by Facility City')
plt.show()

# Correlation matrix
correlation_matrix = df.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



Task 3 Data Preprocessing in Python

a) and b) Data Preprocessing

Handling Missing Values

```
In [22]: # Check for missing values
print(df.isnull().sum())

# If you want to drop rows with any missing values
df.dropna(inplace=True)

# If you want to fill missing values with mean/median
# For example, filling missing values with mean for numerical columns
numerical_cols = df.select_dtypes(include=['int64', 'float64']).columns
for col in numerical_cols:
    df[col].fillna(df[col].mean(), inplace=True)
```

```
Year                                0
Ontario GHG ID                     0
Facility Owner                      0
Facility Name                       0
Facility City                       0
Facility Primary NAICS Code         0
Carbon dioxide (CO2) from non-biomass in CO2e (t) 0
Carbon dioxide (CO2) from biomass in CO2e (t)    0
Methane (CH4) in CO2e (t)           0
Nitrous oxide (N2O) in CO2e (t)     0
Sulphur hexafluoride (SF6) in CO2e (t) 0
Hydrofluorocarbons (HFCs) in CO2e (t) 0
Perfluorocarbons (PFCs) in CO2e (t)  0
Nitrogen Trifluoride (NF3) in CO2e (t) 0
Total CO2e from all sources in CO2e (t) 0
Reporting Amount in CO2e (t)        0
Verification Amount in CO2e (t)     0
Accredited Verification Body         0
dtype: int64
```

Handling Catogorical Data

```
In [18]: # Perform one-hot encoding for categorical columns
categorical_cols = df.select_dtypes(include=['object']).columns
df_encoded = pd.get_dummies(df, columns=categorical_cols, drop_first=True)
```

Feature Scaling

```
In [19]: from sklearn.preprocessing import MinMaxScaler

# Apply Min-Max scaling to numerical columns
scaler = MinMaxScaler()
df_encoded[numerical_cols] = scaler.fit_transform(df_encoded[numerical_cols])
```

Outlier Handling

```
In [20]: # Outlier handling us
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        df.loc[(df[column] < lower_bound) | (df[column] > upper_bound), column] = 0

# Call the function for each numerical column
for col in numerical_cols:
    handle_outliers_iqr(col)
```

Data Splitting

```
In [21]: from sklearn.model_selection import train_test_split

X = df_encoded.drop('Total CO2e from all sources in CO2e (t)', axis=1)
y = df_encoded['Total CO2e from all sources in CO2e (t)']

# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random
```

Task 4 Implementing Missing Values

a) Two Machine learning Model

```
In [23]: from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Initialize the Random Forest Regressor
rf_model = RandomForestRegressor(random_state=42)

# Train the model
rf_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred_rf = rf_model.predict(X_test)

# Evaluate the model
rmse_rf = mean_squared_error(y_test, y_pred_rf, squared=False)
r2_rf = r2_score(y_test, y_pred_rf)

print("Random Forest Model:")
print("Root Mean Squared Error (RMSE):", rmse_rf)
print("R-squared (R2) Score:", r2_rf)
```

```
Random Forest Model:
Root Mean Squared Error (RMSE): 0.00529040819151716
R-squared (R2) Score: 0.9979629664088066
```

```
In [24]: from sklearn.svm import SVR

# Initialize the Support Vector Regressor
svm_model = SVR(kernel='linear')

# Train the model
svm_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred_svm = svm_model.predict(X_test)

# Evaluate the model
rmse_svm = mean_squared_error(y_test, y_pred_svm, squared=False)
r2_svm = r2_score(y_test, y_pred_svm)

print("\nSupport Vector Machine Model:")
print("Root Mean Squared Error (RMSE):", rmse_svm)
print("R-squared (R2) Score:", r2_svm)
```

Support Vector Machine Model:
Root Mean Squared Error (RMSE): 0.08977304389489155
R-squared (R2) Score: 0.413440965899481

b) Evaluate and Compare

```
In [25]: from sklearn.metrics import mean_squared_error, r2_score

# Evaluate Random Forest Model
y_pred_rf = rf_model.predict(X_test)
rmse_rf = mean_squared_error(y_test, y_pred_rf, squared=False)
r2_rf = r2_score(y_test, y_pred_rf)

# Evaluate Support Vector Machine (SVM) Model
y_pred_svm = svm_model.predict(X_test)
rmse_svm = mean_squared_error(y_test, y_pred_svm, squared=False)
r2_svm = r2_score(y_test, y_pred_svm)

# Display results
print("Random Forest Model:")
print("Root Mean Squared Error (RMSE):", rmse_rf)
print("R-squared (R2) Score:", r2_rf)

print("\nSupport Vector Machine Model:")
print("Root Mean Squared Error (RMSE):", rmse_svm)
print("R-squared (R2) Score:", r2_svm)
```

Random Forest Model:

Root Mean Squared Error (RMSE): 0.00529040819151716

R-squared (R2) Score: 0.9979629664088066

Support Vector Machine Model:

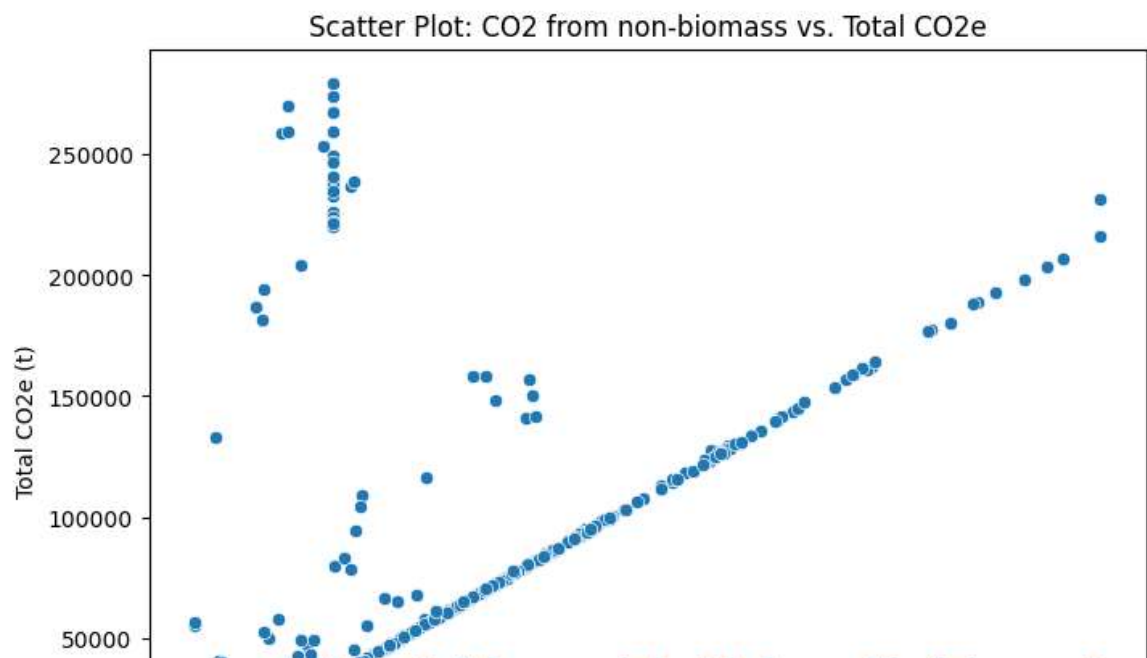
Root Mean Squared Error (RMSE): 0.08977304389489155

R-squared (R2) Score: 0.413440965899481


```
In [26]: # Scatter plot: CO2 from non-biomass vs. Total CO2e
plt.figure(figsize=(8, 6))
sns.scatterplot(x='Carbon dioxide (CO2) from non-biomass in CO2e (t)', y='Total CO2e (t)')
plt.xlabel('CO2 from non-biomass (t)')
plt.ylabel('Total CO2e (t)')
plt.title('Scatter Plot: CO2 from non-biomass vs. Total CO2e')
plt.show()

# Heatmap: Correlation matrix
correlation_matrix = df.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()

# Bar chart: Facility City vs. Total CO2e
plt.figure(figsize=(12, 6))
sns.barplot(x='Facility City', y='Total CO2e from all sources in CO2e (t)', data=df)
plt.xticks(rotation=45, ha='right')
plt.xlabel('Facility City')
plt.ylabel('Total CO2e (t)')
plt.title('Bar Chart: Facility City vs. Total CO2e')
plt.show()
```



```
In [29]: import pip
await pip.install('plotly')
```

```
In [30]: import plotly.express as px
```

```
# 1. Scatter plot: CO2 from non-biomass vs. Total CO2e (with Seaborn)
```

```
plt.figure(figsize=(8, 6))
sns.scatterplot(x='Carbon dioxide (CO2) from non-biomass in CO2e (t)', y='Total CO2e (t)')
plt.xlabel('CO2 from non-biomass (t)')
plt.ylabel('Total CO2e (t)')
plt.title('Scatter Plot: CO2 from non-biomass vs. Total CO2e')
plt.show()
```

```
# 2. Heatmap: Correlation matrix (with Seaborn)
```

```
correlation_matrix = df.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```

```
# 3. Bar chart: Facility City vs. Total CO2e (with Seaborn)
```

```
plt.figure(figsize=(12, 6))
sns.barplot(x='Facility City', y='Total CO2e from all sources in CO2e (t)', data=df)
plt.xticks(rotation=45, ha='right')
plt.xlabel('Facility City')
plt.ylabel('Total CO2e (t)')
plt.title('Bar Chart: Facility City vs. Total CO2e')
plt.show()
```

```
# 4. Interactive bar chart: Facility City vs. Total CO2e (with Plotly)
```

```
fig = px.bar(df, x='Facility City', y='Total CO2e from all sources in CO2e (t)')
fig.update_layout(xaxis_title='Facility City', yaxis_title='Total CO2e (t)')
fig.show()
```

```
# 5. Line chart: Year vs. Total CO2e (with Plotly)
```

```
fig = px.line(df, x='Year', y='Total CO2e from all sources in CO2e (t)', title='Total CO2e (t)')
fig.update_layout(xaxis_title='Year', yaxis_title='Total CO2e (t)')
fig.show()
```

