

Task 2 Data Exploration with python

a) Perform Exploratory data analysis

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("wholesalegasolineprices.csv")

# Overview of the dataset
print(df.head())
print(df.info())
print(df.describe())
```

	Date	Day-of Toronto Wholesale Gasoline	Prior Day NY Harbor Spot \
0	6/25/2018	83.5	NaN
1	6/26/2018	82.8	70.3
2	6/27/2018	83.4	70.9
3	6/28/2018	85.6	72.6
4	6/29/2018	85.3	72.6

	Day-of Toronto less Prior Day NY Harbor \
0	NaN
1	12.5
2	12.5
3	13.0
4	12.7

	Day-of Thunder Bay Wholesale Gasoline	Prior Day Edmonton Spot \
0	0.0	NaN
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

	Day-of Thunder Bay less Prior Day Edmonton
0	NaN
1	0.0
2	0.0
3	0.0
4	0.0

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1297 entries, 0 to 1296
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	Date	1297 non-null	object
1	Day-of Toronto Wholesale Gasoline	1297 non-null	float64
2	Prior Day NY Harbor Spot	1296 non-null	float64
3	Day-of Toronto less Prior Day NY Harbor	1296 non-null	float64
4	Day-of Thunder Bay Wholesale Gasoline	1297 non-null	float64
5	Prior Day Edmonton Spot	1296 non-null	float64
6	Day-of Thunder Bay less Prior Day Edmonton	1296 non-null	float64

```
dtypes: float64(6), object(1)
```

```
memory usage: 65.9+ KB
```

```
None
```

	Day-of Toronto Wholesale Gasoline	Prior Day NY Harbor Spot \
count	1297.000000	1296.000000
mean	77.759445	71.744753
std	22.199763	24.134664
min	23.300000	16.600000
25%	62.700000	56.600000
50%	76.300000	68.900000
75%	91.100000	88.900000
max	146.700000	149.700000

	Day-of Toronto less Prior Day NY Harbor \
count	1296.000000
mean	6.010262
std	3.578425
min	-6.400000

25%	3.600000
50%	6.600000
75%	8.400000
max	16.600000

	Day-of Thunder Bay Wholesale Gasoline	Prior Day Edmonton Spot \
count	1297.000000	1296.000000
mean	61.429453	57.056019
std	44.157387	41.688146
min	0.000000	0.000000
25%	0.000000	0.000000
50%	70.600000	64.450000
75%	95.300000	89.150000
max	155.600000	149.400000

	Day-of Thunder Bay less Prior Day Edmonton
count	1296.000000
mean	4.423765
std	2.805891
min	0.000000
25%	0.000000
50%	6.200000
75%	6.200000
max	8.200000

```


In [3]: # Convert the 'Date' column to datetime format
df['Date'] = pd.to_datetime(df['Date'])

# Plot each column as a separate line chart
plt.figure(figsize=(10, 6))
plt.plot(df['Date'], df['Day-of Toronto Wholesale Gasoline'], label='Day-of Toronto Wholesale Gasoline')
plt.plot(df['Date'], df['Prior Day NY Harbor Spot'], label='Prior Day NY Harbor Spot')
plt.plot(df['Date'], df['Day-of Toronto less Prior Day NY Harbor'], label='Day-of Toronto less Prior Day NY Harbor')
plt.plot(df['Date'], df['Day-of Thunder Bay Wholesale Gasoline'], label='Day-of Thunder Bay Wholesale Gasoline')
plt.plot(df['Date'], df['Prior Day Edmonton Spot'], label='Prior Day Edmonton Spot')
plt.plot(df['Date'], df['Day-of Thunder Bay less Prior Day Edmonton'], label='Day-of Thunder Bay less Prior Day Edmonton')

# Set labels and title
plt.xlabel('Date')
plt.ylabel('Value')
plt.title('Gasoline Prices')
plt.legend()

# Show the plot
plt.tight_layout()
plt.show()

```




b) Generate summary stats

```
In [4]: # Generate summary statistics
summary_stats = df.describe()

# Identify data types
data_types = df.dtypes

# Visualize data distribution
plt.figure(figsize=(10, 6))
for column in df.columns[1:]:
    plt.plot(df['Date'], df[column], label=column)
plt.xlabel('Date')
plt.ylabel('Value')
plt.title('Gasoline Prices')
plt.show()

print("Summary Statistics:")
print(summary_stats)
print("\nData Types:")
print(data_types)
```



Summary Statistics:

	Day-of Toronto Wholesale Gasoline	Prior Day NY Harbor Spot \
count	1297.000000	1296.000000
mean	77.759445	71.744753
std	22.199763	24.134664
min	23.300000	16.600000
25%	62.700000	56.600000
50%	76.300000	68.900000
75%	91.100000	88.900000
max	146.700000	149.700000

	Day-of Toronto less Prior Day NY Harbor \
count	1296.000000
mean	6.010262
std	3.578425
min	-6.400000
25%	3.600000
50%	6.600000
75%	8.400000
max	16.600000

	Day-of Thunder Bay Wholesale Gasoline	Prior Day Edmonton Spot \
count	1297.000000	1296.000000
mean	61.429453	57.056019
std	44.157387	41.688146
min	0.000000	0.000000
25%	0.000000	0.000000
50%	70.600000	64.450000
75%	95.300000	89.150000
max	155.600000	149.400000

	Day-of Thunder Bay less Prior Day Edmonton
count	1296.000000
mean	4.423765
std	2.805891
min	0.000000
25%	0.000000
50%	6.200000
75%	6.200000
max	8.200000

Data Types:

Date	datetime64[ns]
Day-of Toronto Wholesale Gasoline	float64
Prior Day NY Harbor Spot	float64
Day-of Toronto less Prior Day NY Harbor	float64
Day-of Thunder Bay Wholesale Gasoline	float64
Prior Day Edmonton Spot	float64
Day-of Thunder Bay less Prior Day Edmonton	float64
dtype:	object

Task 3 Data preprocessing with python

a) Preprocess

```
In [5]: '''
plt.figure(figsize=(10, 6))
for i, column in enumerate(df.columns[1:], 1):
    plt.subplot(2, 3, i)
    plt.bar(df['Date'], df[column])
    plt.xlabel('Date')
    plt.ylabel(column)

plt.tight_layout()
plt.show()
'''
```

```
Out[5]: "\nplt.figure(figsize=(10, 6))\nfor i, column in enumerate(df.columns[1:], 1):\n    plt.subplot(2, 3, i)\n    plt.bar(df['Date'], df[column])\n    plt.xlabel('Date')\n    plt.ylabel(column)\n\nplt.tight_layout()\nplt.show()\n"
```

b) Handle Missing values, Outliers, and Perform

```
In [6]: # Convert the 'Date' column to datetime format
df['Date'] = pd.to_datetime(df['Date'])

# Set the 'Date' column as the DataFrame index
df.set_index('Date', inplace=True)

df = df.fillna(0)

threshold = df.quantile(0.99)
df = df.clip(upper=threshold, axis=1)

df['Diff_Toronto_and_NY'] = df['Day-of Toronto Wholesale Gasoline'] - df['Prior Day NY Harbor']
df['Diff_Toronto_and_ThunderBay'] = df['Day-of Toronto less Prior Day NY Harbor'] - df['Prior Day Thunder Bay']
df['Diff_Edmonton_and_ThunderBay'] = df['Day-of Thunder Bay less Prior Day Edmonton'] - df['Prior Day Thunder Bay']
```



```
In [7]: df.head(10)
```

Out[7]:

Date	Day-of Toronto Wholesale Gasoline	Prior Day NY Harbor Spot	Day-of Toronto less Prior Day NY Harbor	Day-of Thunder Bay Wholesale Gasoline	Prior Day Edmonton Spot	Day-of Thunder Bay less Prior Day Edmonton	Diff_Toronto_and_NY	Diff_
2018-06-25	83.5	0.0	0.000	0.0	0.0	0.0	83.5	
2018-06-26	82.8	70.3	12.404	0.0	0.0	0.0	12.5	
2018-06-27	83.4	70.9	12.404	0.0	0.0	0.0	12.5	
2018-06-28	85.6	72.6	12.404	0.0	0.0	0.0	13.0	
2018-06-29	85.3	72.6	12.404	0.0	0.0	0.0	12.7	
2018-07-02	85.3	73.6	11.700	0.0	0.0	0.0	11.7	
2018-07-03	86.1	72.4	12.404	0.0	0.0	0.0	13.7	
2018-07-04	84.7	72.7	12.000	0.0	0.0	0.0	12.0	
2018-07-05	84.7	72.6	12.100	0.0	0.0	0.0	12.1	
2018-07-06	84.8	73.1	11.700	0.0	0.0	0.0	11.7	

Task 4 Implement Machine learning models with python

a) Two machine learning models

```
In [8]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Split the data into features (X) and target (y)
X = df.drop(columns=['Day-of Toronto Wholesale Gasoline'])
y = df['Day-of Toronto Wholesale Gasoline']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Model 1: Linear Regression
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
lr_predictions = lr_model.predict(X_test)

# Model 2: Random Forest Regression
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
rf_predictions = rf_model.predict(X_test)

# Evaluate Model Performance
lr_mse = mean_squared_error(y_test, lr_predictions)
lr_r2 = r2_score(y_test, lr_predictions)

rf_mse = mean_squared_error(y_test, rf_predictions)
rf_r2 = r2_score(y_test, rf_predictions)

# Compare the performance of the models
print("Linear Regression:")
print(f"Mean Squared Error: {lr_mse:.2f}")
print(f"R-squared: {lr_r2:.2f}")

print("\nRandom Forest Regression:")
print(f"Mean Squared Error: {rf_mse:.2f}")
print(f"R-squared: {rf_r2:.2f}")
```

Linear Regression:

Mean Squared Error: 0.00

R-squared: 1.00

Random Forest Regression:

Mean Squared Error: 1.17

R-squared: 1.00

```
In [9]: import pandas as pd
        from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import confusion_matrix, classification_report

        X = df.drop(columns=['Day-of Toronto Wholesale Gasoline'])
        y = df['Day-of Toronto Wholesale Gasoline']

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random

        model = LogisticRegression()
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)

        cm = confusion_matrix(y_test, y_pred)
        print("Confusion Matrix:")
        print(cm)

        cr = classification_report(y_test, y_pred)
        print("\nClassification Report:")
        print(cr)
```

```

-----
ValueError                                Traceback (most recent call last)
Cell In[9], line 12
      9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=
0.2, random_state=42)
     11 model = LogisticRegression()
--> 12 model.fit(X_train, y_train)
     13 y_pred = model.predict(X_test)
     15 cm = confusion_matrix(y_test, y_pred)

File /lib/python3.11/site-packages/sklearn/linear_model/_logistic.py:1204, in
LogisticRegression.fit(self, X, y, sample_weight)
     1194     _dtype = [np.float64, np.float32]
     1196 X, y = self._validate_data(
     1197     X,
     1198     y,
     (... )
     1202     accept_large_sparse=solver not in ["liblinear", "sag", "saga"],
     1203 )
-> 1204 check_classification_targets(y)
     1205 self.classes_ = np.unique(y)
     1207 multi_class = _check_multi_class(self.multi_class, solver, len(self.c
lasses_))

File /lib/python3.11/site-packages/sklearn/utils/multiclass.py:218, in check_
classification_targets(y)
     210 y_type = type_of_target(y, input_name="y")
     211 if y_type not in [
     212     "binary",
     213     "multiclass",
     (... )
     216     "multilabel-sequences",
     217 ]:
--> 218     raise ValueError("Unknown label type: %r" % y_type)


ValueError: Unknown label type: 'continuous'

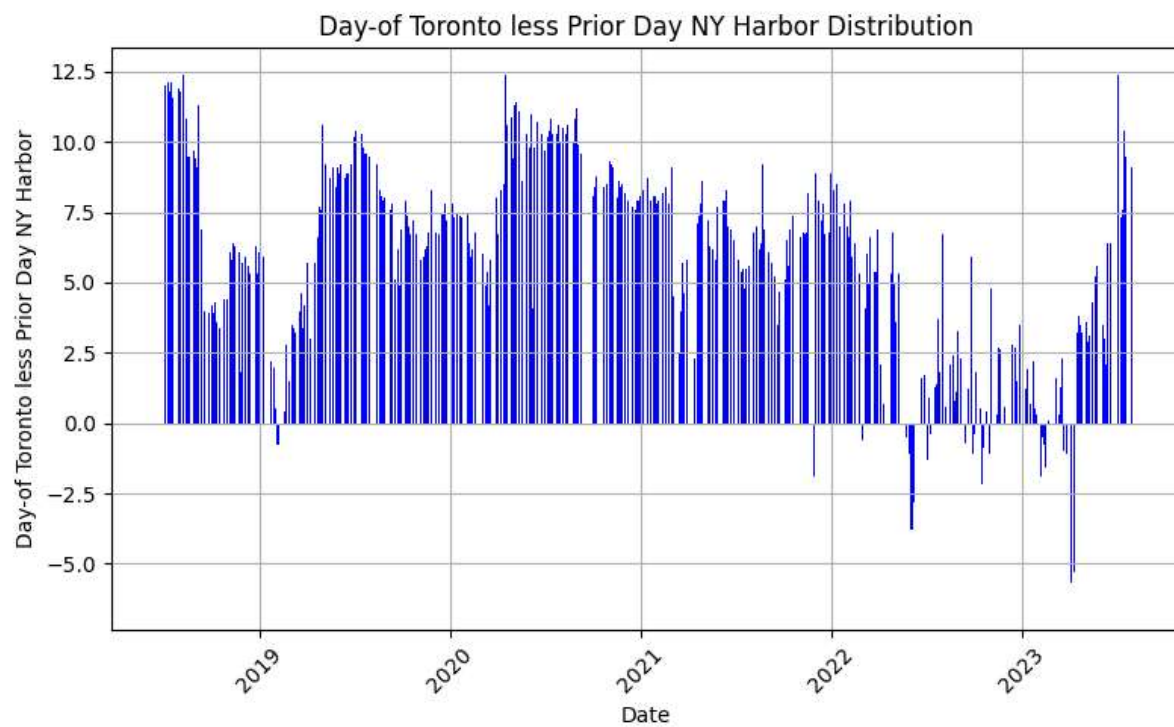
```

```
In [10]: import pandas as pd
import matplotlib.pyplot as plt

# Line Plot: Day-of Toronto Wholesale Gasoline
plt.figure(figsize=(10, 6))
plt.plot(df.index, df['Day-of Toronto Wholesale Gasoline'], marker='o', linestyle='solid')
plt.xlabel('Date')
plt.ylabel('Day-of Toronto Wholesale Gasoline')
plt.title('Day-of Toronto Wholesale Gasoline Trend')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()

# Bar Chart: Day-of Toronto Less Prior Day NY Harbor
plt.figure(figsize=(8, 5))
plt.bar(df.index, df['Day-of Toronto less Prior Day NY Harbor'], color='blue')
plt.xlabel('Date')
plt.ylabel('Day-of Toronto less Prior Day NY Harbor')
plt.title('Day-of Toronto less Prior Day NY Harbor Distribution')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```





b) Evaluate two models

```
In [11]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Preprocessing (if necessary)

# Split the data into features (X) and target (y)
X = df.drop(columns=['Day-of Toronto Wholesale Gasoline'])
y = df['Day-of Toronto Wholesale Gasoline']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Model 1: Linear Regression
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
lr_predictions = lr_model.predict(X_test)

# Model 2: Random Forest Regression
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
rf_predictions = rf_model.predict(X_test)

# Evaluate Model Performance
lr_mse = mean_squared_error(y_test, lr_predictions)
lr_r2 = r2_score(y_test, lr_predictions)

rf_mse = mean_squared_error(y_test, rf_predictions)
rf_r2 = r2_score(y_test, rf_predictions)

# Compare the performance of the models
print("Linear Regression:")
print(f"Mean Squared Error: {lr_mse:.2f}")
print(f"R-squared: {lr_r2:.2f}")

print("\nRandom Forest Regression:")
print(f"Mean Squared Error: {rf_mse:.2f}")
print(f"R-squared: {rf_r2:.2f}")
```

Linear Regression:
Mean Squared Error: 0.00
R-squared: 1.00

Random Forest Regression:
Mean Squared Error: 1.17
R-squared: 1.00

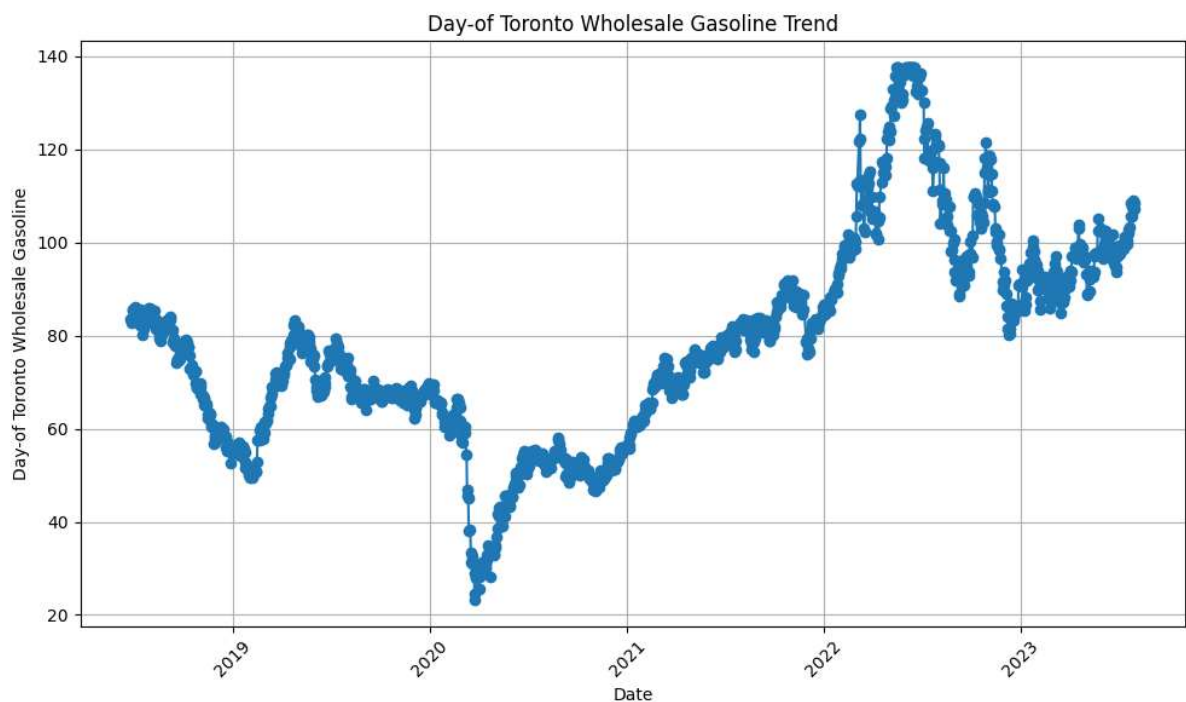
Task 5 Data Visualization

5.1 Create Meaningfull plots


```
In [15]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Line Plot: Day-of Toronto Wholesale Gasoline
plt.figure(figsize=(10, 6))
plt.plot(df.index, df['Day-of Toronto Wholesale Gasoline'], marker='o', linestyle='solid')
plt.xlabel('Date')
plt.ylabel('Day-of Toronto Wholesale Gasoline')
plt.title('Day-of Toronto Wholesale Gasoline Trend')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()

# Bar Chart: Day-of Toronto Less Prior Day NY Harbor
plt.figure(figsize=(8, 5))
sns.barplot(x=df.index, y=df['Day-of Toronto less Prior Day NY Harbor'], color='blue')
plt.xlabel('Date')
plt.ylabel('Day-of Toronto less Prior Day NY Harbor')
plt.title('Day-of Toronto less Prior Day NY Harbor Distribution')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```



Day-of Toronto less Prior Day NY Harbor Distribution



```

In [17]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Sample data
data = {
    'Date': ['6/25/2018', '6/26/2018', '6/27/2018', '6/28/2018', '6/29/2018'],
    'Day-of Toronto Wholesale Gasoline': [83.5, 82.8, 82.9, 83.1, 83.4],
    'Prior Day NY Harbor Spot': [0, 70.3, 72.1, 70.5, 69.8],
    'Day-of Toronto less Prior Day NY Harbor': [0, 12.5, 10.8, 12.6, 13.6],
    'Day-of Thunder Bay Wholesale Gasoline': [0, 0, 0, 0, 0],
    'Prior Day Edmonton Spot': [0, 0, 0, 0, 0],
    'Day-of Thunder Bay less Prior Day Edmonton': [0, 0, 0, 0, 0]
}

# Create DataFrame from the data
df = pd.DataFrame(data)

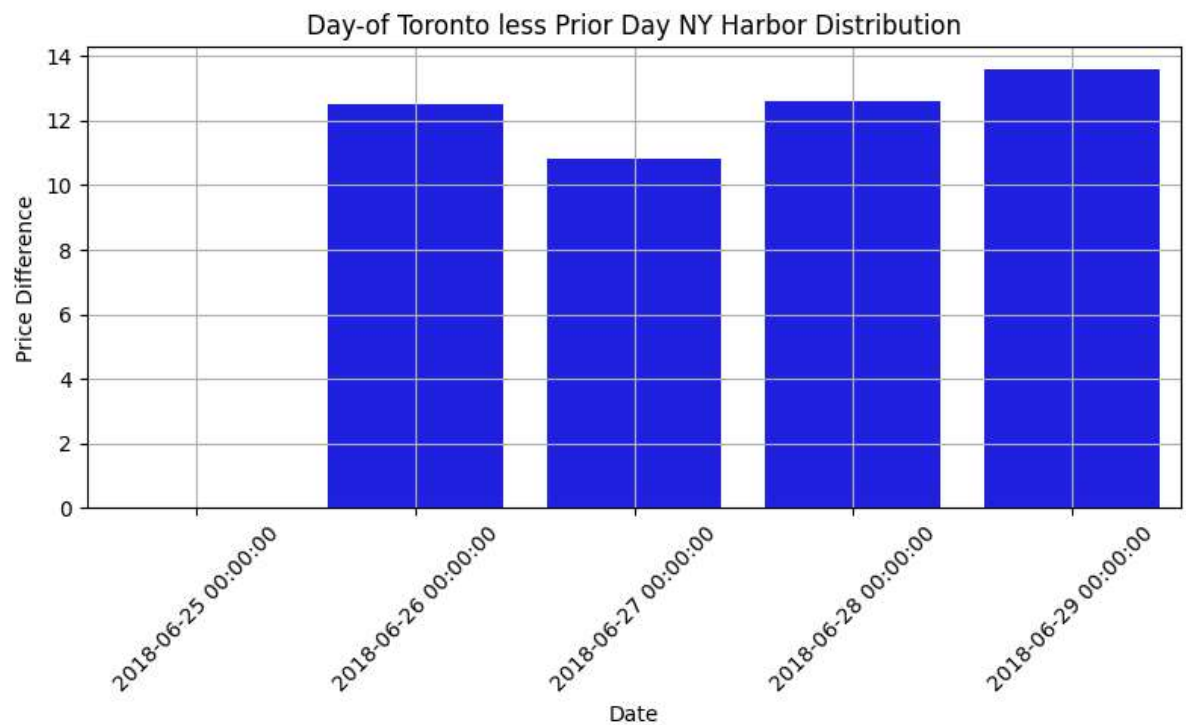

# Convert the 'Date' column to datetime format
df['Date'] = pd.to_datetime(df['Date'])

# Set the 'Date' column as the DataFrame index
df.set_index('Date', inplace=True)

# Line Plot: Day-of Toronto Wholesale Gasoline
plt.figure(figsize=(10, 6))
plt.plot(df.index, df['Day-of Toronto Wholesale Gasoline'], marker='o', linestyle='solid')
plt.xlabel('Date')
plt.ylabel('Gasoline Price')
plt.title('Day-of Toronto Wholesale Gasoline Trend')
plt.xticks(rotation=45)
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# Bar Chart: Day-of Toronto Less Prior Day NY Harbor
plt.figure(figsize=(8, 5))
sns.barplot(x=df.index, y=df['Day-of Toronto less Prior Day NY Harbor'], color='red')
plt.xlabel('Date')
plt.ylabel('Price Difference')
plt.title('Day-of Toronto less Prior Day NY Harbor Distribution')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()

```



In []:

In []: